

Navigating the Multilingual Landscape of Scientific Computing:

Python, Julia, and Awkward Array

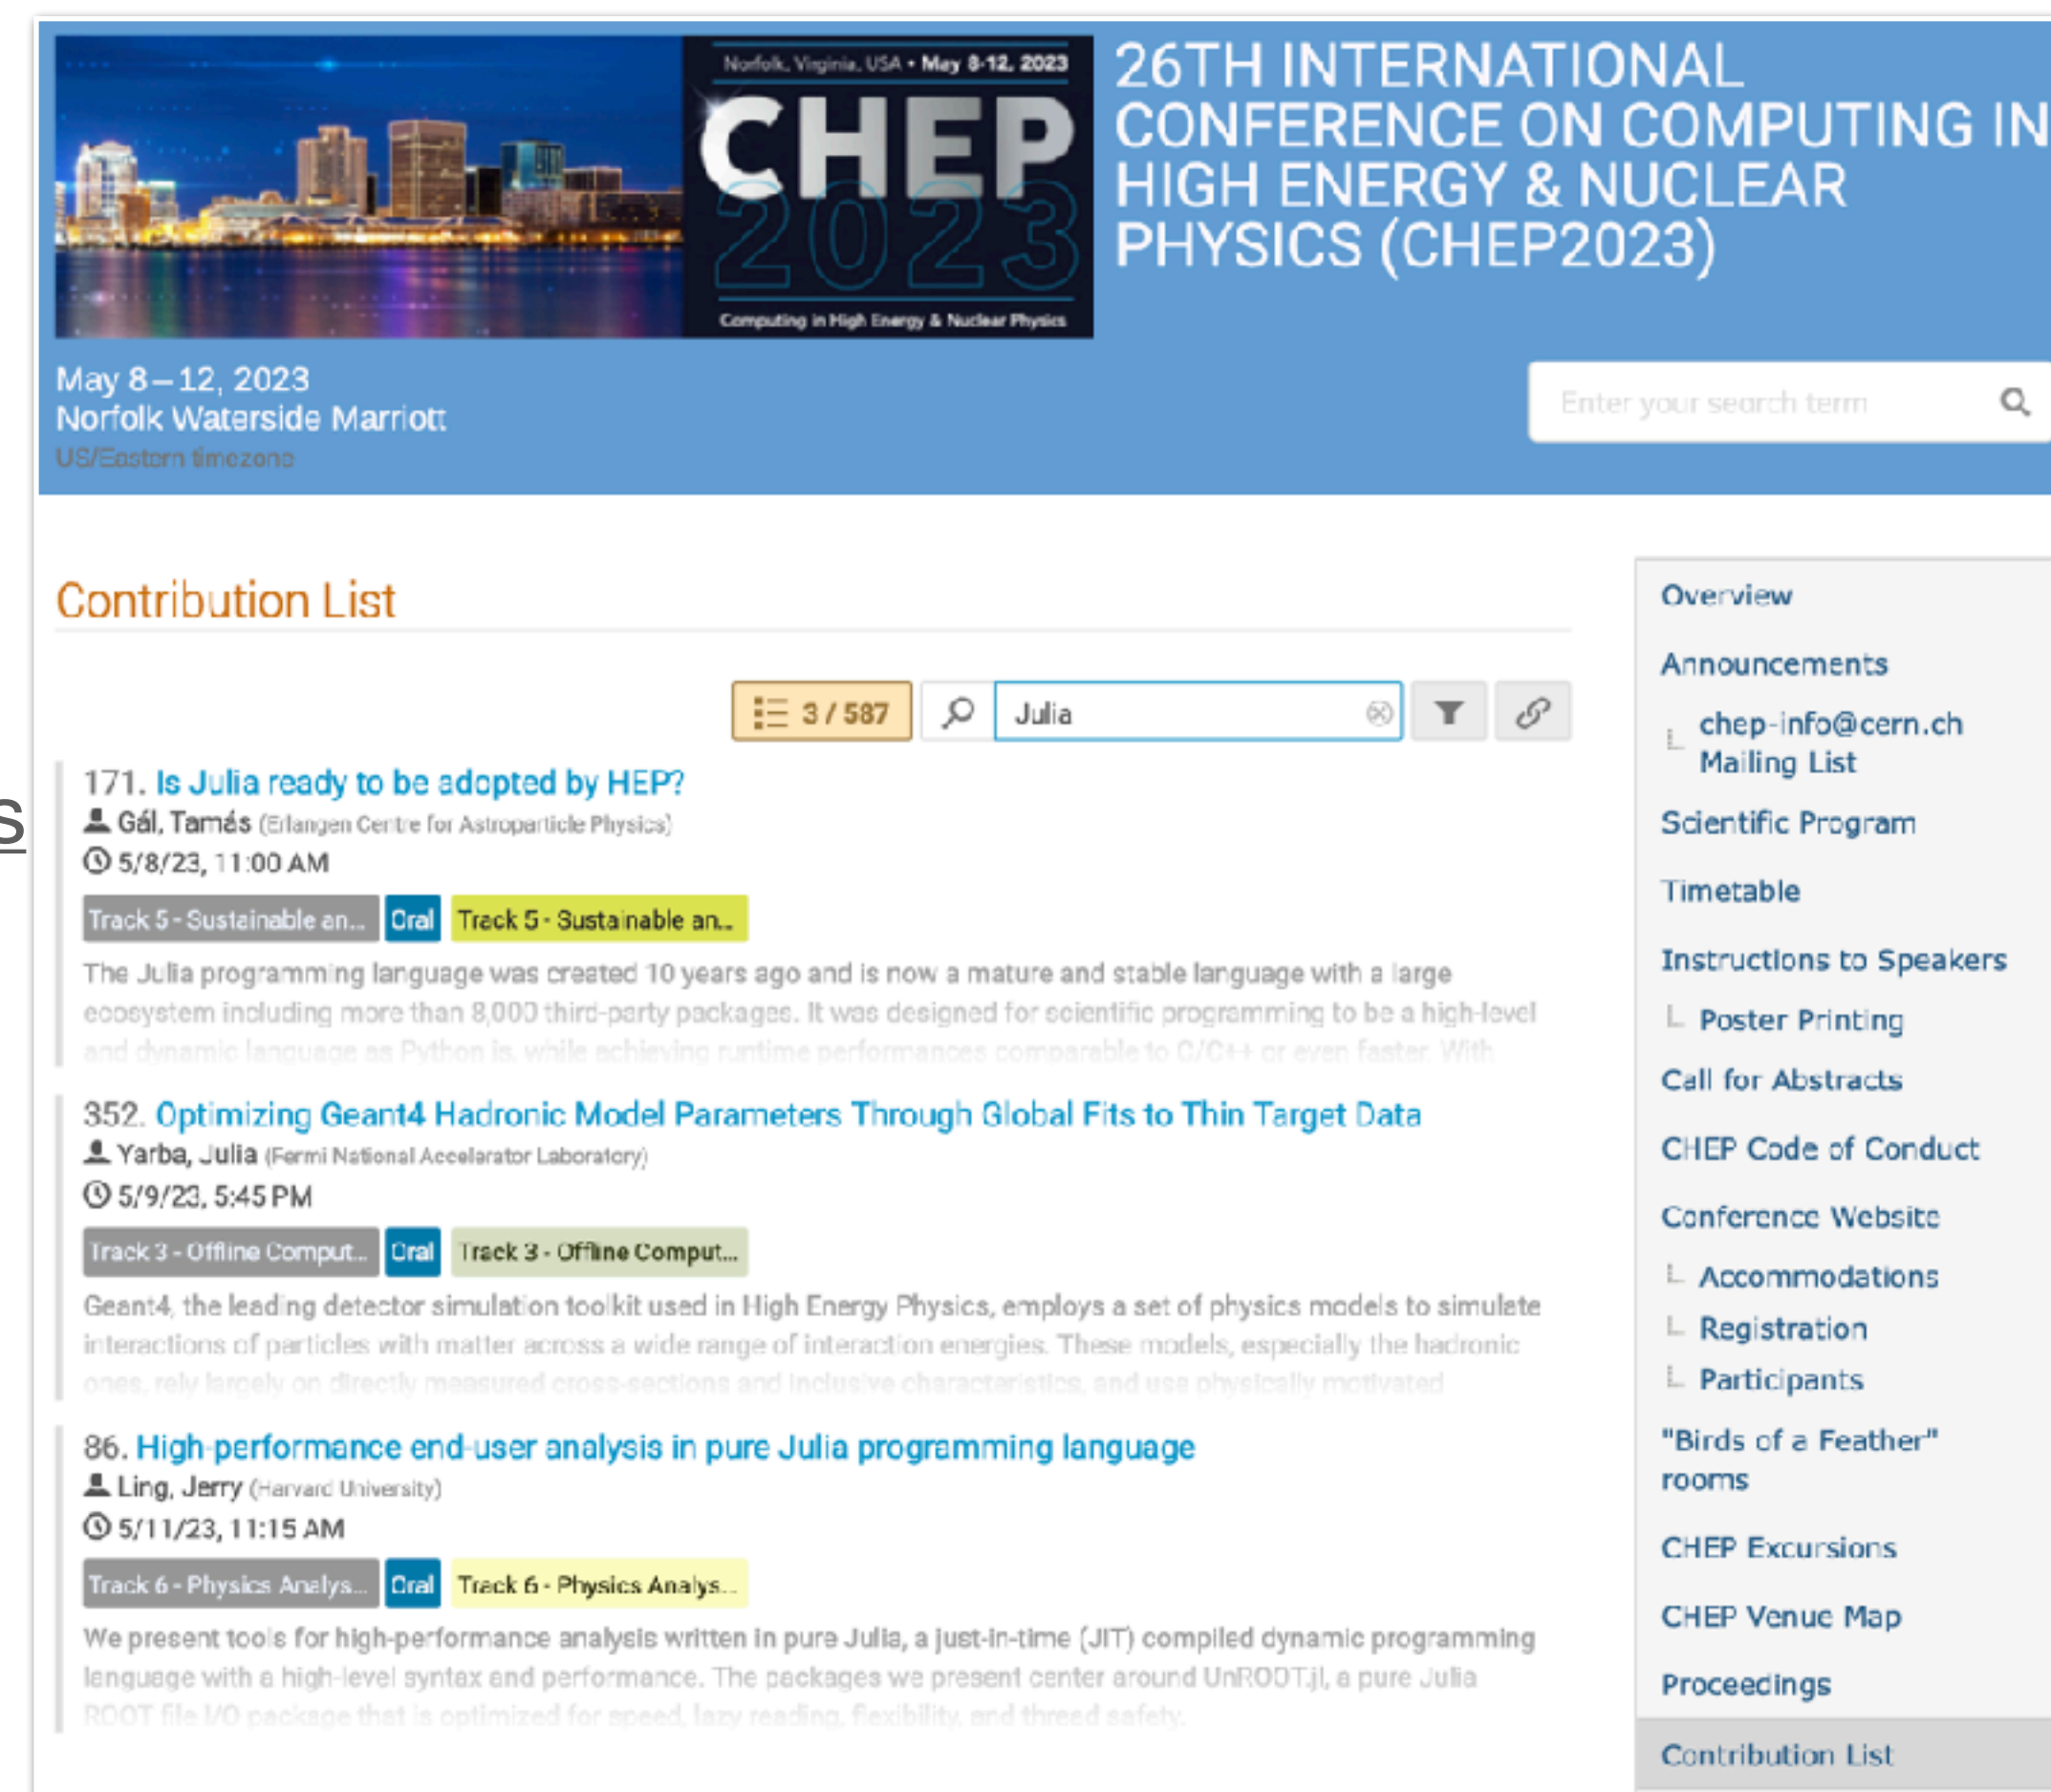
Ianna Osborne, Jim Pivarski, Jerry 🦄 Ling



Julia at CHEP2024 and CHEP2023

a new trend?

- [Julia in HEP](#) by Graeme A Stewart, 21 Oct 2024, [Plenary session](#)
- [R&D towards heterogenous frameworks for Future Experiments](#) by Mateusz Jakub Fila, 21 Oct 2024, [Parallel \(Track 3\)](#)
- [ROOT RNTuple implementation in Julia programming language](#) by Jerry Ling, 21 Oct 2024, [Parallel \(Track 5\)](#)
- [Comparative efficiency of HEP codes across languages and architectures](#) by Samuel Cadellin Skipsey, 21 Oct 2024, [Parallel \(Track 6\)](#)
- [EDM4hep.jl: Analysing EDM4hep files with Julia](#) by Pere Mato, 21 Oct 2024, [Poster session](#)
- [Fast Jet Reconstruction in Julia](#) by Graeme A Stewart, 23 Oct 2024, [Parallel \(Track 3\)](#)
- [BAT.jl, the Bayesian Analysis Toolkit in Julia](#) by Oliver Schulz, 23 Oct 2024, [Parallel \(Track 5\)](#)
- [Navigating the Multilingual Landscape of Scientific Computing: Python, Julia, and Awkward Array](#), by Ianna Osborne, 24 Oct 2024, [Parallel \(Track 9\)](#)



The screenshot shows the CHEP 2023 website. The header includes the conference title "26TH INTERNATIONAL CONFERENCE ON COMPUTING IN HIGH ENERGY & NUCLEAR PHYSICS (CHEP2023)" and the dates "May 8-12, 2023" at the "Norfolk Waterside Marriott". A search bar is present. The "Contribution List" section shows search results for "Julia". Three results are visible:

- 171. [Is Julia ready to be adopted by HEP?](#) by Gál, Tamás (Erlangen Centre for Astroparticle Physics), 5/8/23, 11:00 AM. Track 5 - Sustainable an... Oral.
- 352. [Optimizing Geant4 Hadronic Model Parameters Through Global Fits to Thin Target Data](#) by Yarba, Julia (Fermi National Accelerator Laboratory), 5/9/23, 5:45 PM. Track 3 - Offline Comput... Oral.
- 86. [High-performance end-user analysis in pure Julia programming language](#) by Ling, Jerry (Harvard University), 5/11/23, 11:15 AM. Track 6 - Physics Analys... Oral.

A right-hand sidebar contains navigation links: Overview, Announcements, chep-info@cern.ch Mailing List, Scientific Program, Timetable, Instructions to Speakers, Poster Printing, Call for Abstracts, CHEP Code of Conduct, Conference Website, Accommodations, Registration, Participants, "Birds of a Feather" rooms, CHEP Excursions, CHEP Venue Map, Proceedings, and Contribution List.

Embedding Julia in Python

How easy is it to blend these languages?

- We can use PythonCall for integrating Python's vast ecosystem into Julia projects and JuliaCall for embedding high-performance Julia code into Python scripts.

```
[ ]: from juliacall import Main as jl
```

```
[ ]: %load_ext juliacall
```

```
[ ]: %%julia

using Pkg
Pkg.add("UnROOT")
using UnROOT
```



Using Julia Packages from Python

```
[1]: from juliacall import Main as jl
```

Detected IPython. Loading juliacall extension. See <https://juliapy.github.io/PythonCall.jl/stable/compat/#IPython>

```
[2]: %load_ext juliacall
```

WARNING: replacing module _ipython.

```
[3]: %%julia
```

```
using Pkg  
Pkg.add("UnROOT")  
using UnROOT
```



Resolving package versions...
No Changes to `~/anaconda3/envs/julia_hep_2024/julia_env/Project.toml`
No Changes to `~/anaconda3/envs/julia_hep_2024/julia_env/Manifest.toml`

ROOT File as Julia Object in Python

Using UnROOT

- This dataset contains about 60 mio. data events from the CMS detector taken in 2012 during Run B and C. The original AOD dataset is converted to the NanoAOD format and reduced to the muon collections.
- Wunsch, Stefan; (2019). DoubleMuParked dataset from 2012 in NanoAOD format reduced on muons. CERN Open Data Portal. DOI:[10.7483/OPENDATA.CMS.LVG5.QT81](https://doi.org/10.7483/OPENDATA.CMS.LVG5.QT81)

```
[175]: %%julia

using UnROOT

@time big_tree = ROOTFile("../.../Run2012BC_DoubleMuParked_Muons.root")

0.007673 seconds (4.65 k allocations: 10.119 MiB)
```

```
[175]: ROOTFile with 2 entries and 17 streamers.
.../.../Run2012BC_DoubleMuParked_Muons.root
└─ Events (TTree)
   └─ "nMuon"
      └─ "Muon_pt"
         └─ "Muon_eta"
            └─ "Muon_phi"
               └─ "Muon_mass"
                  └─ "Muon_charge"
```

```
[174]: %%julia

@time events = LazyTree(big_tree, "Events")

0.000334 seconds (365 allocations: 31.703 KiB)
```

```
[174]: 61,540,413 rows × 6 columns (omitted printing of 61,540,403 rows)
```

	Muon_phi	nMuon	Muon_pt	Muon_eta	Muon_charge	Muon_mass
	SubArray{Float3	UInt32	SubArray{Float3	SubArray{Float3	SubArray{Int32,	SubArray{Float3
1	[-0.0343, 2.54]	2	[10.8, 15.7]	[1.07, -0.564]	[-1, -1]	[0.106, 0.106]
2	[-0.275, 2.54]	2	[10.5, 16.3]	[-0.428, 0.349]	[1, -1]	[0.106, 0.106]
3	[-1.22]	1	[3.28]	[2.21]	[1]	[0.106]
4	[-2.08, 0.251, -2.01, -1.85]	4	[11.4, 17.6, 9.62, 3.5]	[-1.59, -1.75, -1.59, -1.66]	[1, 1, 1, 1]	[0.106, 0.106, 0.106, 0.106]

Julia ROOT Tree in Python

Faster way to read ROOT files

```
[7]: events = jl.Main.LazyTree(file, "Events")
```

```
[8]: %%timeit  
jl.Main.LazyTree(file, "Events")
```

368 μ s \pm 23.2 μ s per loop (mean \pm std. dev. of 7 runs, 1,000 loops each)

- With viewing the data as AwkwardArray we can use either Julia or Python analysis code or even combine both languages.
- Getting the best performance from Julia requires us to focus on type-stability and good practices for reducing unnecessary recompilation.

Including Julia Code in Python


Notes on code organization

```

53 # Predefine the output structure with a concrete NamedTuple type
54 const RecordArrayType = NamedTuple{(:pt, :eta, :phi, :mass, :charge, :isolation)}
55
56 function make_record_array(
57     events::NamedTuple{(:muon,)},
58     Tuple{
59         NamedTuple{(:pt, :eta, :phi, :mass, :charge, :pfRelIso03_all),
60             Tuple{
61                 Vector{T}, Vector{T}, Vector{T}, Vector{T}, Vector{T}, Vector{T}
62             }
63         }
64     }
65 ) where T
66
67     # Convert the relevant fields into AwkwardArray arrays
68     array = AwkwardArray.RecordArray(
69         RecordArrayType((
70             AwkwardArray.from_iter(events.muon.pt),
71             AwkwardArray.from_iter(events.muon.eta),
72             AwkwardArray.from_iter(events.muon.phi),
73             AwkwardArray.from_iter(events.muon.mass),
74             AwkwardArray.from_iter(events.muon.charge),
75             AwkwardArray.from_iter(events.muon.pfRelIso03_all)
76         ))
77     )
78
79     return array
80 end

```

```
[20]: jl.include('awkward_analyzer_functions.jl');
```




- Provide the correct path when using the include function.
- If your project grows larger, consider structuring your code into more modules and files for better organization:
 - It is generally a good practice to organize your code into modules. This helps with namespace management and reduces the likelihood of name collisions.
 - Use export to expose functions from a module. This makes it easy to access the desired functionality after including a module.

AwkwardArray.jl as Data Bridge

Between Julia and Python

- Using AwkwardArray in Julia to calculate Higgs mass:



for Julia!

Getting started

Let's assume that both Python and Julia are installed.

Note

If Julia is not install it is recomend to follow its official installation instructions described [here](#).

Installation

It is recommended to use [conda](#) virtul environment.

Using Julia Awkward Arrays from Python

- To install [Awkward Array](#) Python package:

```
conda install -c conda-forge awkward
```

- To install [JuliaCall](#):

```
conda install pyjuliacall
```

JuliaCall takes care of installing all necessary Julia packages, including this package.

```
import awkward as ak
from juliacall import Main as jl

jl.seval("using AwkwardArray")
```

Getting started

Introduction

Getting started

- Installation
- Using Julia Awkward Arrays from Python
- Using Python Awkward Arrays from Julia

Converting Arrays

API

Types

Functions

LICENSE

Version dev

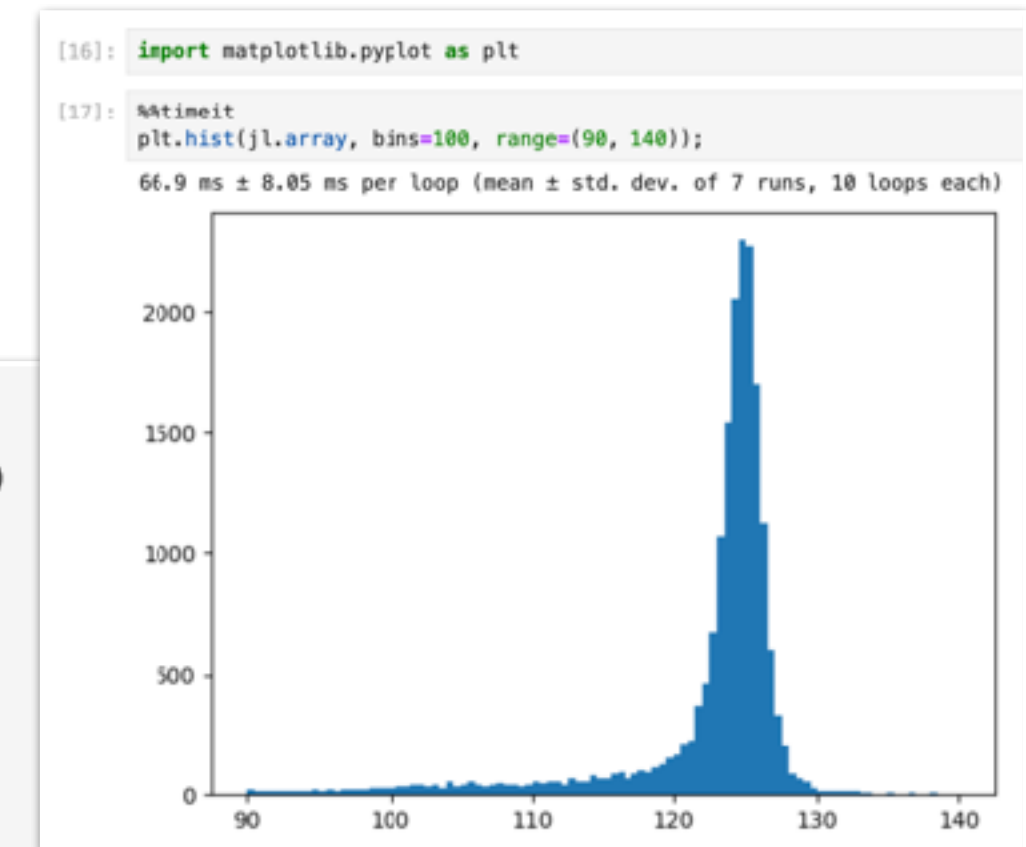
```
function main_looper(events)
    array = AwkwardArray.PrimitiveArray{Float64}()
    for evt in events

        (; Muon_charge) = evt
        if length(Muon_charge) != 4
            continue
        end
        sum(Muon_charge) != 0 && continue # shortcut if-else

        (; Muon_pt, Muon_eta, Muon_phi, Muon_mass) = evt
        higgs_4vector = sum(LorentzVectorCyl.(Muon_pt, Muon_eta, Muon_phi, Muon_mass))
        higgs_mass = mass(higgs_4vector)

        push!(array, higgs_mass)
    end

    return array
end
```



Calling Julia from Python Efficiency

with a very small overhead

	Call Julia from Python	Julia	Julia precompiled
Open ROOT file with UnROOT.ROOTFile	1.05 ms \pm 28.5 μ s per loop (mean \pm std. dev. of 7 runs, 1,000 loops each)	0.047574 seconds (6.00 k allocations: 565.602 KiB, 94.32% compilation time)	0.001293 seconds (4.78 k allocations: 509.227 KiB)
Get a tree with UnROOT.LazyTree	368 μ s \pm 23.2 μ s per loop (mean \pm std. dev. of 7 runs, 1,000 loops each)	0.520122 seconds (468.07 k allocations: 32.848 MiB, 99.26% compilation time)	0.000502 seconds (2.40 k allocations: 230.648 KiB)
Execute Julia function main_looper	CPU times: user 430 ms, sys: 22 ms, total: 452 ms Wall time: 452 ms	0.236601 seconds (27.00 k allocations: 120.572 MiB, 6.44% compilation time)	0.473383 seconds (17.49 k allocations: 120.002 MiB, 45.32% gc time, 6.97% compilation time: 100% of which was recompilation)
			0.226617 seconds (3.08 k allocations: 118.958 MiB, 5.16% gc time)

?

AwkwardArray.jl Overhead

Compared with Using Typed Arrays

```
[12]: %%julia

array = @time main_looper(events)

0.459656 seconds (398.55 k allocations: 77.552 MiB, 2.98% gc time, 75.63% compilation time)
```

```
[12]: 20525-element AwkwardArray.PrimitiveArray{Float64, Vector{Float64}, :default}:
125.12303161621094
123.90653991699219
124.15757751464844
122.6549301147461
125.26071166992188
124.77593994140625
124.20553588867188
124.42249298095703
110.03680419921875
124.46846008300781
⋮
127.15644836425781
70.50875091552734
```

```
[30]: %%julia

array = @time main_looper(events)

0.225840 seconds (185.55 k allocations: 73.431 MiB, 2.74% gc time, 46.84% compilation time)
```

```
[30]: 20525-element Vector{Float64}:
125.12303161621094
123.90653991699219
124.15757751464844
122.6549301147461
125.26071166992188
124.77593994140625
124.20553588867188
124.42249298095703
110.03680419921875
124.46846008300781
⋮
```

```
[38]: %%julia


array = @time main_looper(events)

0.257670 seconds (185.09 k allocations: 71.296 MiB, 1.61% gc time, 45.21% compilation time)

[38]: 20525-element AwkwardArray.PrimitiveArray{Float64, Vector{Float64}, :default}:
125.12303161621094
123.90653991699219
124.15757751464844
122.6549301147461
125.26071166992188
124.77593994140625
124.20553588867188
124.42249298095703
110.03680419921875
124.46846008300781
⋮
```

- Started with using AwkwardArray and compared it to a Julia native typed array: Vector

- Takeaway:** no significant overhead seen after small changes to Julia `main_looper` code.



But can we do better?

Small Code Changes in Destructure and Skip


```
function main_looper(events)
    array = AwkwardArray.PrimitiveArray{Float64}()
    for evt in events

        (; Muon_charge) = evt
        if length(Muon_charge) != 4
            continue
        end
        sum(Muon_charge) != 0 && continue # shortcut if-else

        (; Muon_pt, Muon_eta, Muon_phi, Muon_mass) = evt
        higgs_4vector = sum(LorentzVectorCyl.(Muon_pt, Muon_eta, Muon_phi, Muon_mass))
        higgs_mass = mass(higgs_4vector)

        push!(array, higgs_mass)
    end

    return array
end
```

 ChatGPT

```
function main_looper(events)
    # Create an empty AwkwardArray for storing the Higgs mass values
    array = AwkwardArray.PrimitiveArray{Float64}()

    # Loop over events and process only valid ones
    for evt in events
        # Destructure the necessary fields from the event
        (; Muon_charge, Muon_pt, Muon_eta, Muon_phi, Muon_mass) = evt

        # Skip event if it doesn't meet the required conditions
        if length(Muon_charge) != 4 || sum(Muon_charge) != 0
            continue
        end

        # Create Lorentz vectors for the muons and calculate the Higgs mass
        higgs_4vector = sum(LorentzVectorCyl.(Muon_pt, Muon_eta, Muon_phi, Muon_mass))
        higgs_mass = mass(higgs_4vector)

        # Add the result to the AwkwardArray
        push!(array, higgs_mass)
    end

    # Return the final AwkwardArray containing Higgs masses
    return array
end
```

- Time reduced from 0.459656 seconds to 0.145763 seconds

But can we do better?

```
[52]: %julia

array = @time main_looper(events)
0.145763 seconds (22.94 k allocations: 59.588 MiB, 2.19% gc time)
[52]: 20525-element AwkwardArray.PrimitiveArray{Float64, Vector{Float64}, :default}:
125.12303161621094
123.90653991699219
124.15757751464844
122.6549301147461
125.26071166992188
```

Ensuring Type Stability

```
function main_looper(events::AwkwardArray.RecordArray)
    # Pre-allocate an AwkwardArray to store Higgs mass values
    array = AwkwardArray.PrimitiveArray{Float64}(undef, length(events))
    count = 0 # To track valid entries

    for evt in events
        # Destructure the necessary fields from the event with concrete types
        (; Muon_charge::Vector{Float64}, Muon_pt::Vector{Float64}, Muon_eta::Vector{Float64},
         Muon_phi::Vector{Float64}, Muon_mass::Vector{Float64}) = evt

        # Check conditions with inlined logic
        if length(Muon_charge) != 4 || sum(Muon_charge) != 0
            continue
        end

        # Compute the Lorentz vector sum and Higgs mass
        higgs_4vector = sum(LorentzVectorCyl.(Muon_pt, Muon_eta, Muon_phi, Muon_mass))
        higgs_mass = mass(higgs_4vector)

        # Store the result in the pre-allocated array
        count += 1
        array[count] = higgs_mass
    end

    # Resize the array to only include valid entries
    return AwkwardArray.subarray(array, 1:count)
end
```



ChatGPT

- Avoid unnecessary recompilation!

But can we do better?

```
[56]: %%julia

array = @time main_looper(events)
0.119881 seconds (22.94 k allocations: 59.588 MiB, 3.07% gc time)
[56]: 20525-element AwkwardArray.PrimitiveArray{Float64, Vector{Float64}, :default}:
125.12303161621094
123.90653991699219
124.15757751464844
122.6549301147461
125.26071166992188
124.77593994140625
124.20553588867188
124.42249298095703
110.03680419921875
124.46846008300781
```


Multithreading Support

Experimental in JuliaCall

```
% export JULIA_NUM_THREADS=4
```

```
% export PYTHON_JULIACALL_HANDLE_SIGNALS=yes
```

```
[89]: %%julia

using AwkwardArray
using Base.Threads

function main_looper_awkward(events)
    array = AwkwardArray.PrimitiveArray{Float64}()
    lock_obj = ReentrantLock() # Create a lock object to control access to shared array

    @threads for i in 1:length(events)
        evt = events[i]


        # Destructure the necessary fields from the event
        (; Muon_charge, Muon_pt, Muon_eta, Muon_phi, Muon_mass) = evt

        # Skip event if it doesn't meet the required conditions
        if length(Muon_charge) != 4 || sum(Muon_charge) != 0
            continue
        end

        # Create Lorentz vectors for the muons and calculate the Higgs mass
        higgs_4vector = sum(LorentzVectorCyl.(Muon_pt, Muon_eta, Muon_phi, Muon_mass))
        higgs_mass = mass(higgs_4vector)

        # Use lock to safely push! into the shared array
        lock(lock_obj) # Explicitly lock before modifying shared data
        try
            push!(array, higgs_mass)
        finally
            unlock(lock_obj) # Ensure the lock is always released
        end
    end

    return array
end
```



- **Execution time reduced by 88%**, speeding up from 0.5 seconds to 0.06 seconds—a **8.33x performance improvement**.
- **Memory usage optimized**, cutting allocations from 398k to 24k, making the process much more efficient.
- Overall, the code is **much faster and leaner**, showing significant gains in both speed and memory management.

```
[91]: %%julia

array = @time main_looper_awkward(events)

0.063811 seconds (24.00 k allocations: 87.739 MiB)

[91]: 20525-element AwkwardArray.PrimitiveArray{Float64, Vector{Float64}, :default}:
 125.12303161621094
 123.90653991699219
 124.15757751464844
```

Performance Scaling

Increasing # Events


```
[111]: %%julia
@time main_looper(events[1:10000])

0.054722 seconds (213.33 k allocations: 65.210 MiB)
[111]: 696-element AwkwardArray.PrimitiveArray{Float64, Vector{
125.12303161621094
123.90653991699219
124.15757751464844
125.54163360595703
123.40087890625
100.83844757080078
102.18019104003906
124.19779968261719
124.63880920410156
124.16439819335938
⋮
```

```
[113]: %%julia
@time main_looper(events[1:100000])

0.216400 seconds (2.11 M allocations: 305.127 MiB)
[113]: 6785-element AwkwardArray.PrimitiveArray{Float64, Vector{
126.94151306152344
126.7888412475586
125.12303161621094
125.17405700683594
121.99496459960938
123.90653991699219
126.61412811279297
124.15757751464844
105.9798583984375
125.23162078857422
⋮
```

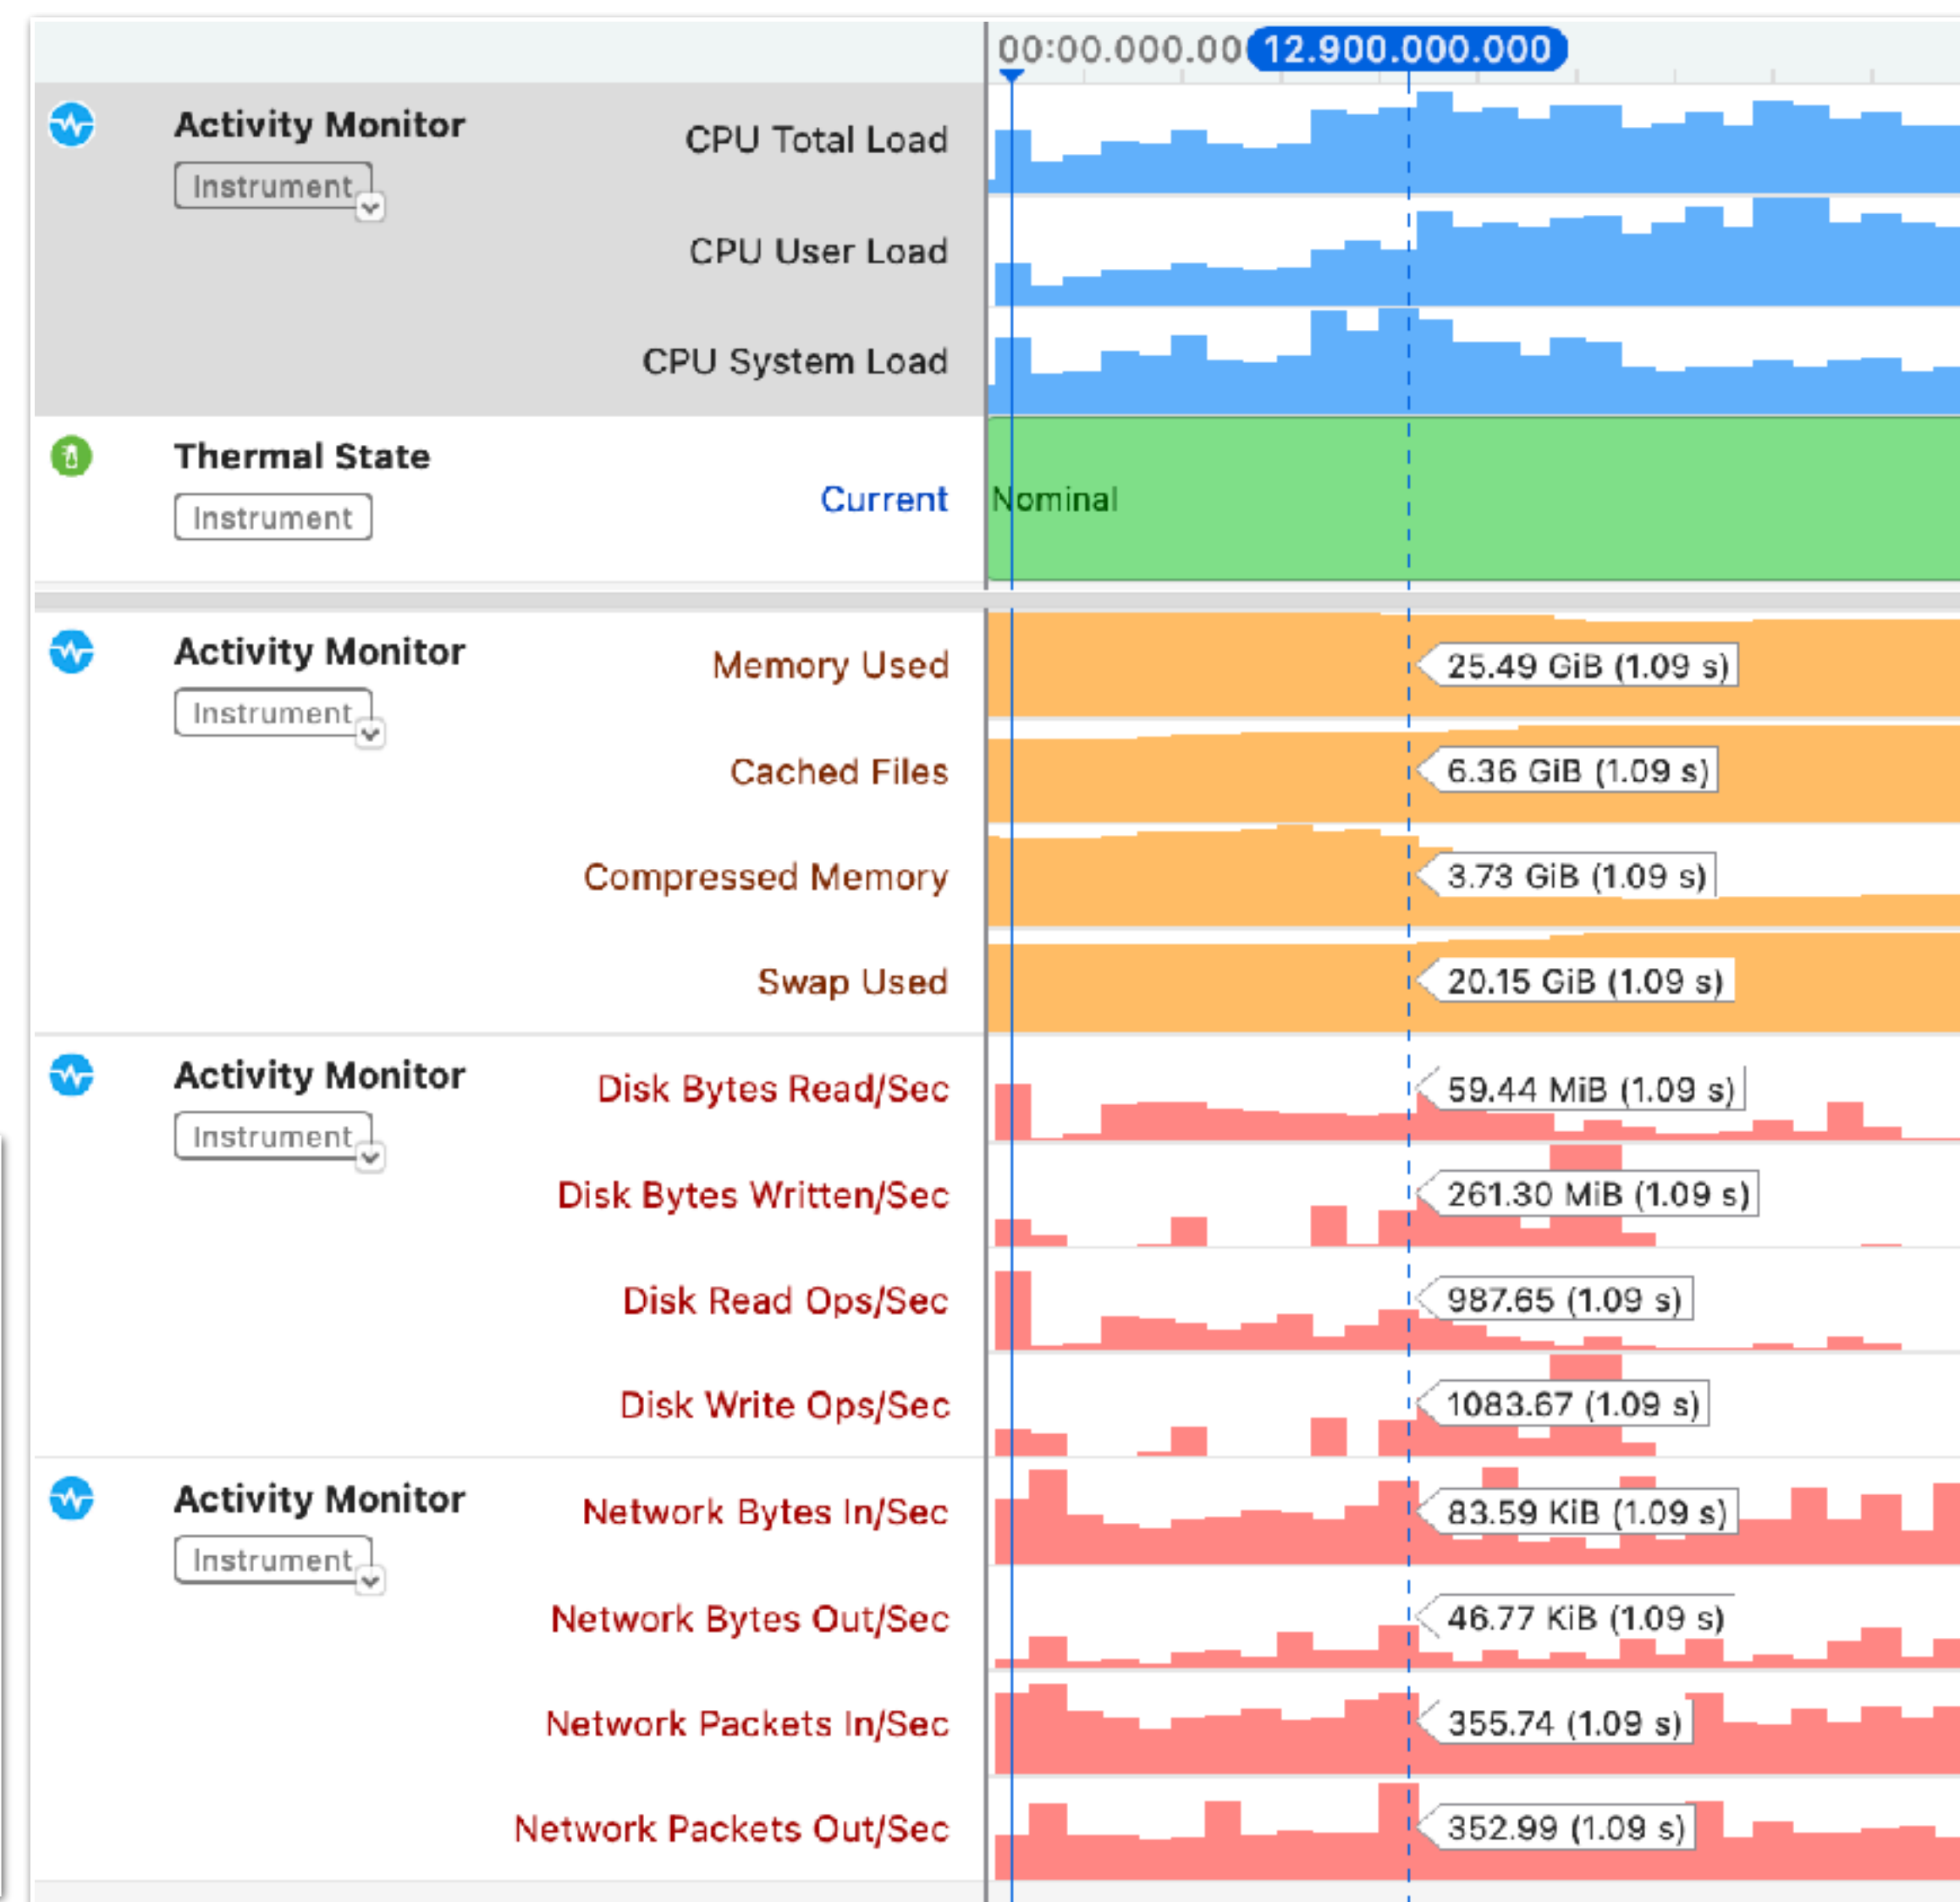
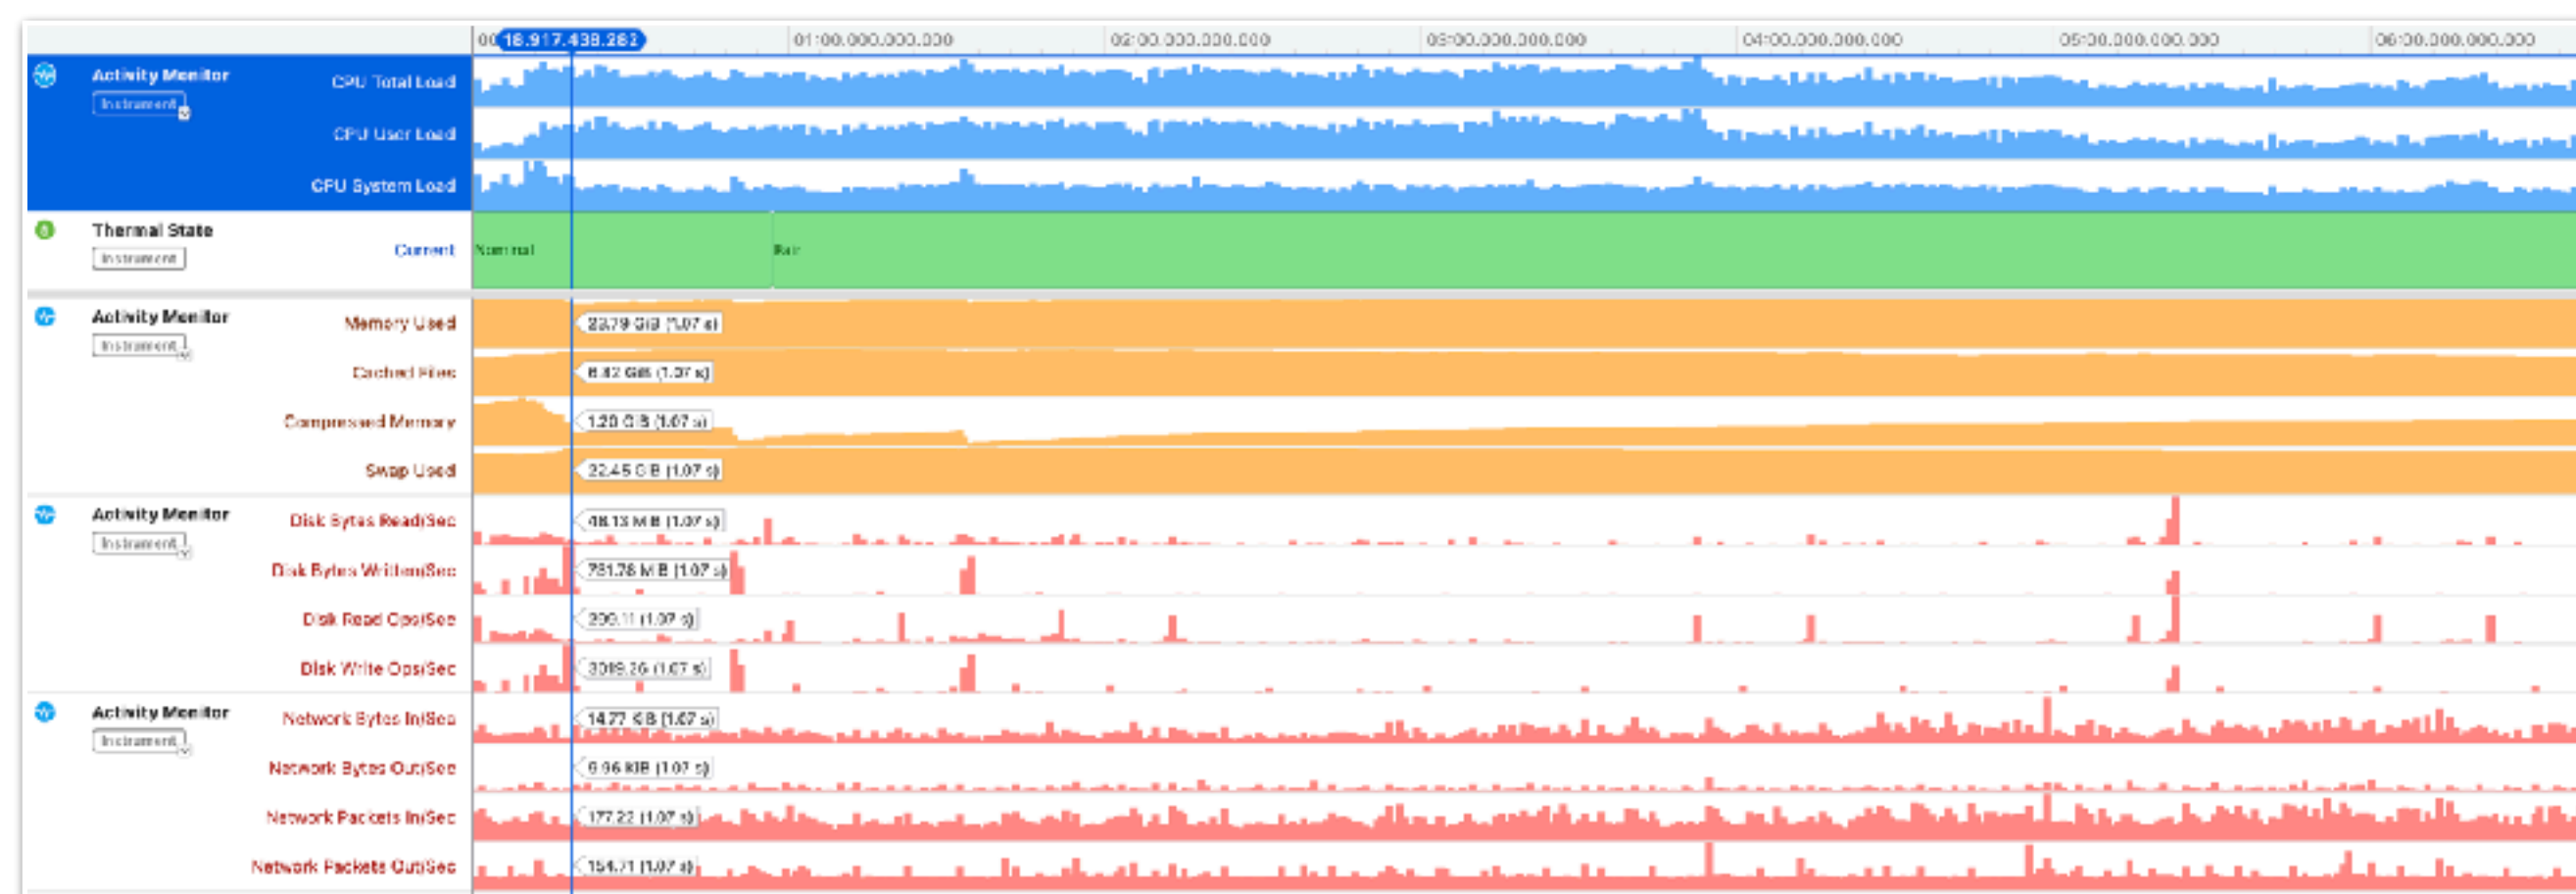
```
@time main_looper(events)
61,540,413
9.897901 seconds (2.09 M allocations: 15.559 GiB, 16.83% gc time)
1915776-element AwkwardArray.PrimitiveArray{Float64, Vector{Float64}, :default}:
```

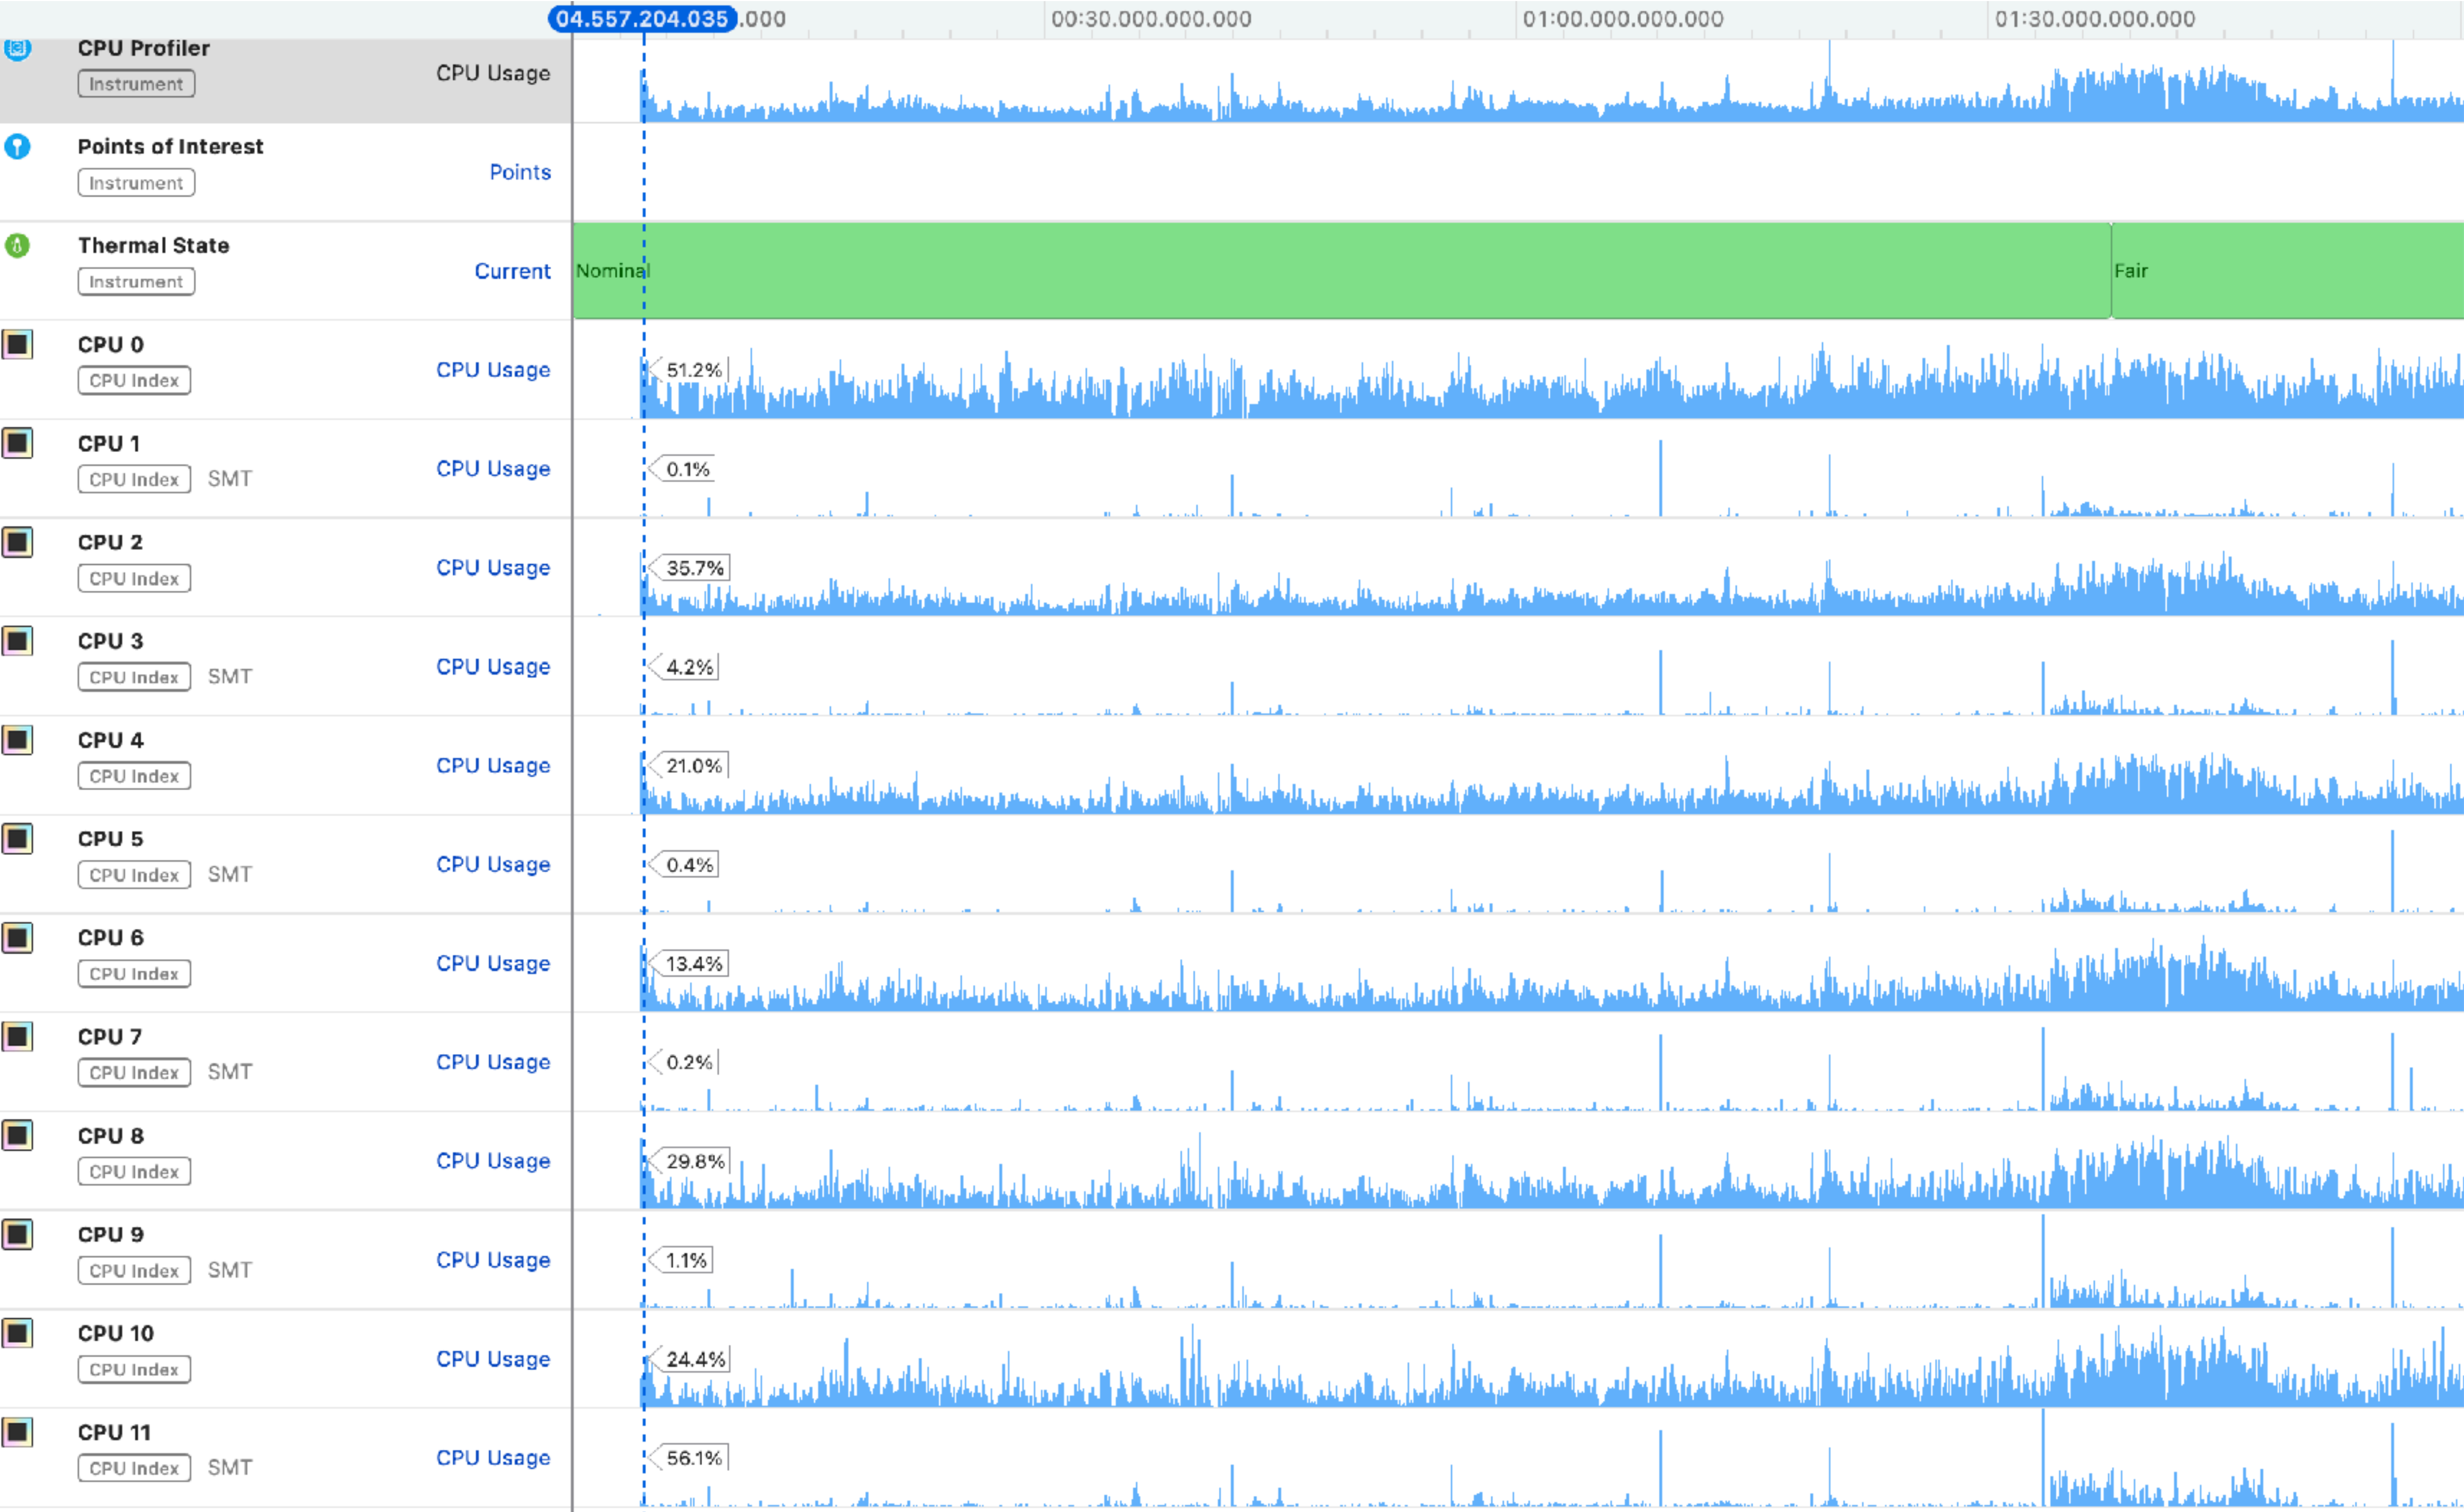


Instruments

Activity Monitor

- macOS Big Sur version 11.6
- Processor 2.6 GHz 6-Core Intel Core i7
- Memory 32 GB 2667 MHz DDR4





But can we do better?

Multi-processing and Distributed Computing

```
[7]: %%julia

@everywhere include("DummyAwkwardModule.jl")

[8]: %%julia

using .DummyAwkwardModule

[9]: %%julia

MuMu = @spawnat :any invariant_mass(events);

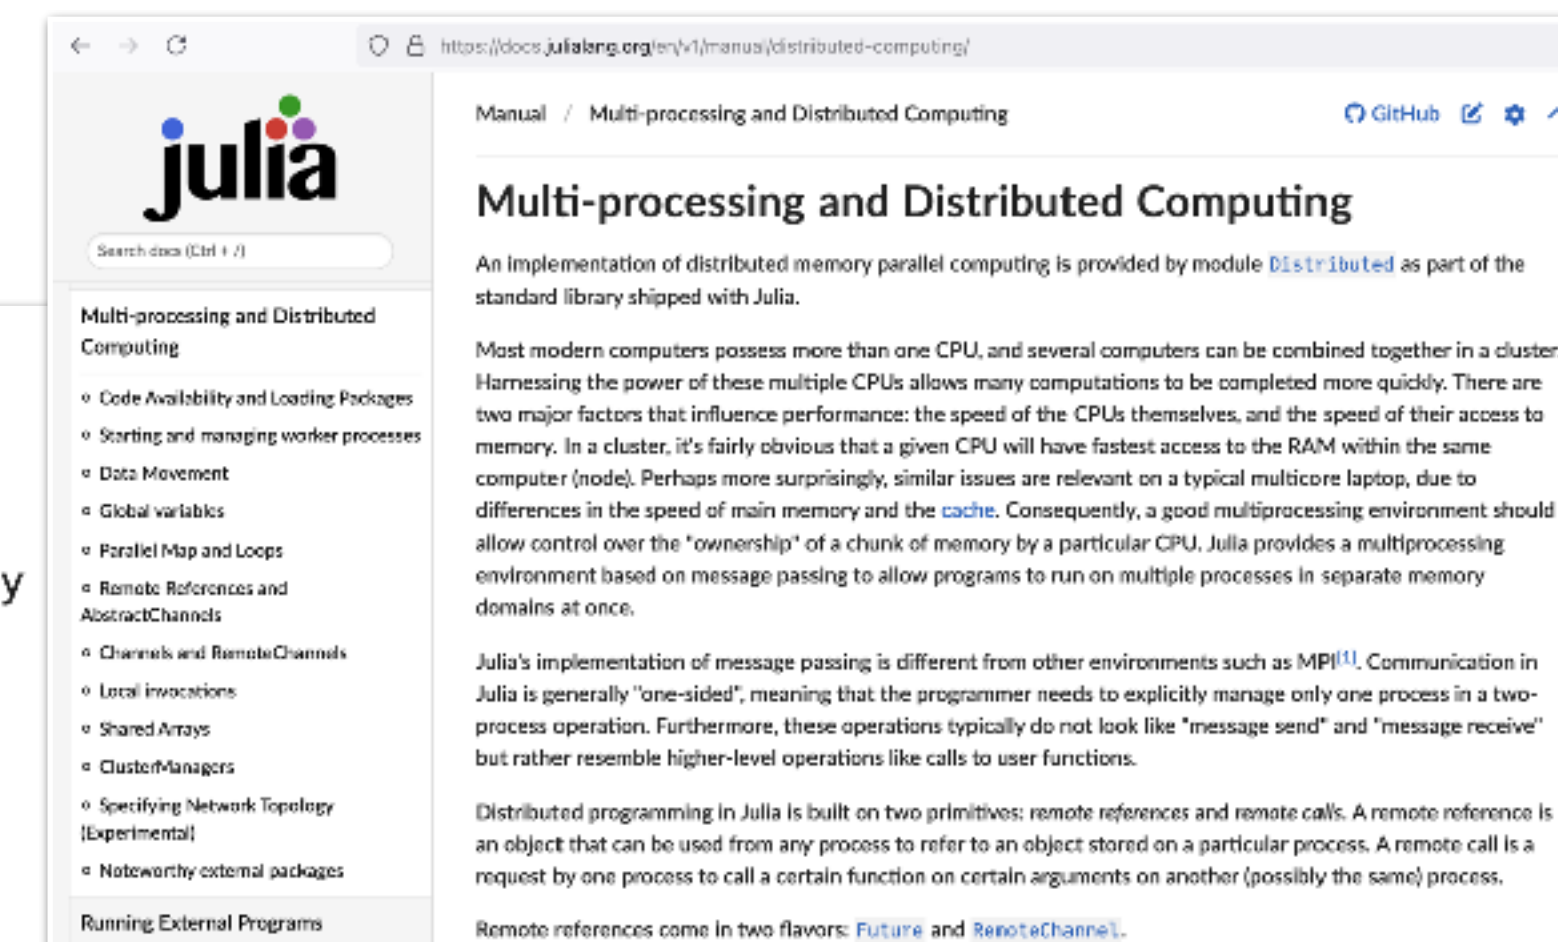
[11]: %%julia

@time fetch(MuMu)

0.000005 seconds

[11]: 5638149-element AwkwardArray.PrimitiveArray{Float64, Vector{Float64}}, :default
113.64686584472656
88.29710388183594
88.33483123779297
91.27149963378906
93.55725860595703
90.91211700439453
89.15238952636719
82.29732513427734
94.57678985595703
89.23975372314453
⋮
```

```
1 module DummyAwkwardModule
2
3 export invariant_mass
4
5 using LorentzVectorHEP, AwkwardArray
6 using UnROOT
7
8 using Base.Threads
9
10 function invariant_mass(cms_events)
11
12     array = AwkwardArray.PrimitiveArray{Float64}()
13     lock_obj = ReentrantLock() # Create a lock object to control access to shared array
14
15     @threads for i in 1:length(cms_events)
16         evt = cms_events[i]
17
18         # Destructure the necessary fields from the event
19         (; Muon_charge, Muon_pt, Muon_eta, Muon_phi, Muon_mass, nMuon) = evt
20
21         # Skip event if it doesn't meet the required conditions
22         if nMuon != 2 || Muon_charge[1] == Muon_charge[2]
23             continue
24         end
25
26         # Calculate invariant mass using LorentzVectorHEP for clarity and accuracy
27         muon1 = LorentzVectorCyl(Muon_pt[1], Muon_eta[1], Muon_phi[1], Muon_mass[1])
28         muon2 = LorentzVectorCyl(Muon_pt[2], Muon_eta[2], Muon_phi[2], Muon_mass[2])
29         result = mass(muon1 + muon2)
30
31         # Only add masses greater than 70 GeV
32         if result > 70
33             # Use lock to safely push! into the shared array
34             lock(lock_obj) # Explicitly lock before modifying shared data
35             try
```



Summary

Optimizing Performance with Julia

- While we may not see a significant speedup from replacing NumPy, Awkward, or Numba with Julia in vectorized operations, identifying tasks that don't fit well with these libraries can unlock Julia's true potential.
- Developing custom kernels for specific problems may lead to innovative solutions, even if it's not immediately obvious.
- Despite challenges in multilingual runtime environments and experimental thread support, the ongoing evolution of Julia offers exciting opportunities for performance enhancement.